

CLASSIFICATION USING CONVOLUTIONAL NEURAL NETWORK.

N.SHIVRAM (RA2011026010237)

G.MEGHANATH (RA2011026010219)

A.S.THARUN (RA2011026010217)

PROBLEM STATEMENT:

- The CIFAR-10 dataset consists of 60000 32x32 color images in 10 classes, with 6000 images per class. There are 50000 training images and 10000 test images.
- The dataset is divided into five training batches and one test batch, each with 10000 images. The test batch contains exactly 1000 randomly-selected images from each class. The training batches contain the remaining images in random order, but some training batches may contain more images from one class than another. Between them, the training batches contain exactly 5000 images from each class.
- This project aims to classify the given dataset into 10 unique classes through the use of Convolutional neural network

airplane



automobile



bird



cat



deer



dog



frog



horse



ship



truck



THE NEURAL NETWORK:


- A Convolutional Neural Network is used in this project to classify the dataset
- A Convolutional Neural Network (ConvNet/CNN) is a Deep Learning algorithm that can take in an input image, assign importance (learnable weights and biases) to various aspects/objects in the image, and be able to differentiate one from the other.
- Since our aim is to classify the dataset into 10 different classes, CNN is the best available neural network for the task.

ADVANTAGES OF CNN:

- CNNs eliminate the need for manual feature extraction—the features are learned directly by the CNN.
- CNNs produce highly accurate recognition results.
- CNNs can be retrained for new recognition tasks, enabling you to build on pre-existing networks.

WORKING OF A CNN:

- A convolutional neural network can have tens or hundreds of layers that each learn to detect different features of an image. Filters are applied to each training image at different resolutions, and the output of each convolved image is used as the input to the next layer. The filters can start as very simple features, such as brightness and edges, and increase in complexity to features that uniquely define the object.



- Like other neural networks, a CNN is composed of an input layer, an output layer, and many hidden layers in between.
- These layers perform operations that alter the data with the intent of learning features specific to the data. Three of the most common layers are: convolution, activation or ReLU, and pooling.

CONVOULUTION LAYER:

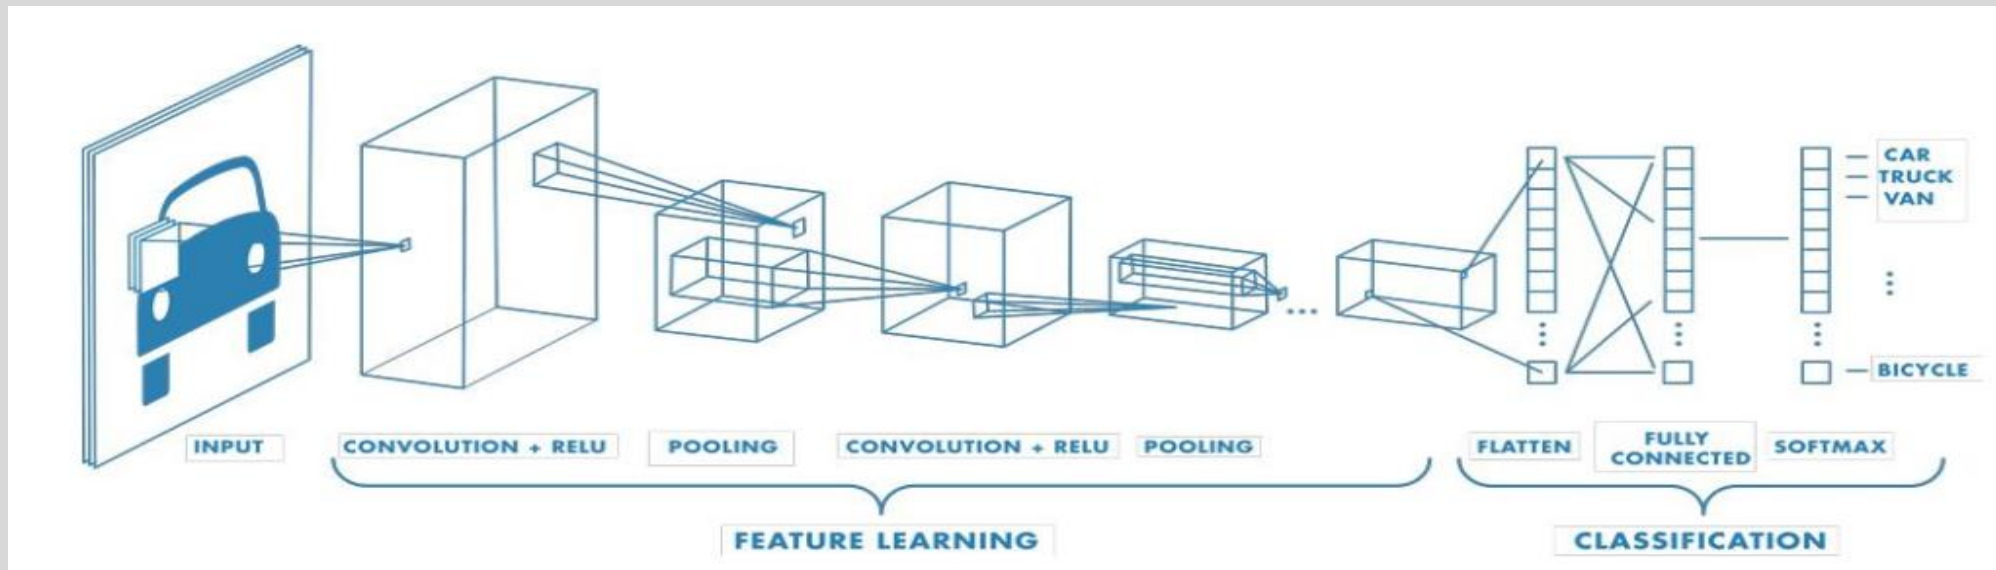
Convolution puts the input images through a set of convolutional filters, each of which activates certain features from the images.


RECTIFIED LINEAR UNIT (ReLU):

ReLU layer allows for faster and more effective training by mapping negative values to zero and maintaining positive values. This is sometimes referred to as activation, because only the activated features are carried forward into the next layer.

POOLING:

Pooling simplifies the output by performing nonlinear down sampling, reducing the number of parameters that the network needs to learn.





DATA FLOW:

- THE CNN HAS 2 sets of convulsion and pooling layers
- A 32x32 image is given as input to the first convolution layer.
- This convolution layer gives a 32 size filter map.
- This 32 size feature map is reduced to 2x2 by the pooling layer.
- This 2x2 feature map is again given to the second convolution layer.
- This layer gives a 64x64 feature map, which is again reduced to 2x2 feature map.
- This 2x2 map is given to the input or the flattening layer, which flattens the inputs.
- The flatten layer transfer the data to dense or hidden layer with 64 neurons.
- After data processing the hidden layer transfers the data to the output layers which have 10 neurons.

MODEL SUMMARY:

... Model: "sequential"


Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 30, 30, 32)	896
max_pooling2d (MaxPooling2D)	(None, 15, 15, 32)	0
conv2d_1 (Conv2D)	(None, 13, 13, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 6, 6, 64)	0
flatten (Flatten)	(None, 2304)	0
dense (Dense)	(None, 64)	147520
dense_1 (Dense)	(None, 10)	650

=====


Total params: 167,562
Trainable params: 167,562
Non-trainable params: 0

=====

ACCURACY AND LOSS GRAPHS



ACCURACY



LOSS

conv2d_input	input:	[(None, 32, 32, 3)]
InputLayer	output:	[(None, 32, 32, 3)]



conv2d	input:	(None, 32, 32, 3)
Conv2D	output:	(None, 30, 30, 32)



max_pooling2d	input:	(None, 30, 30, 32)
MaxPooling2D	output:	(None, 15, 15, 32)



conv2d_1	input:	(None, 15, 15, 32)
Conv2D	output:	(None, 13, 13, 64)



max_pooling2d_1	input:	(None, 13, 13, 64)
MaxPooling2D	output:	(None, 6, 6, 64)



flatten	input:	(None, 6, 6, 64)
Flatten	output:	(None, 2304)



dense	input:	(None, 2304)
Dense	output:	(None, 64)



dense_1	input:	(None, 64)
Dense	output:	(None, 10)

CODE:

```
1 import matplotlib.pyplot as plt
2 import numpy as np
3 import tensorflow as tf
4 from tensorflow.keras import datasets, layers, models
```

[2] ✓ 23.6s

Python

```
1 (x_train, y_train), (x_test, y_test) = datasets.cifar10.load_data()
2 x_train.shape
```

[3] ✓ 1.3s

Python

... (50000, 32, 32, 3)

DATA PROCESSING

```
1 x_test.shape
```

[4] ✓ 0.5s

Python

... (10000, 32, 32, 3)

```
1 y_train.shape
```

[5] ✓ 0.8s

Python

... (50000, 1)


```
1 y_train[:5]
```

[6] ✓ 0.8s

Python

```
... array([[6],  
          [9],  
          [9],  
          [4],  
          [1]], dtype=uint8)
```

```
1 y_train = y_train.reshape(-1,)  
2 y_train[:5]
```

[7] ✓ 0.6s

Python

```
... array([6, 9, 9, 4, 1], dtype=uint8)
```

▷ ▾
1 y_test = y_test.reshape(-1,)

[8] ✓ 0.3s

Python

+ Code

+ Markdown

```
1 classes = ["airplane", "automobile", "bird", "cat", "deer", "dog", "frog", "horse", "ship", "truck"]
```

[9] ✓ 0.4s

Python

```
1 def plot_sample(x, y, index):  
2     plt.figure(figsize = (15,2))  
3     plt.imshow(x[index])  
4     plt.xlabel(classes[y[index]])
```

[10] ✓ 0.4s

Python




```
1 plot_sample(x_train, y_train, 5)
```

[11] ✓ 0.7s

Python

...



```
1 x_train = x_train / 255.0
```

```
2 x_test = x_test / 255.0
```

[12] ✓ 2.4s

Python

CNN



```
1 cnn = models.Sequential([
2     layers.Conv2D(filters=32, kernel_size=(3, 3), activation='relu', input_shape=(32, 32, 3)),
3     layers.MaxPooling2D((2, 2)),
4
5     layers.Conv2D(filters=64, kernel_size=(3, 3), activation='relu'),
6     layers.MaxPooling2D((2, 2)),
7
8     layers.Flatten(),
9     layers.Dense(64, activation='relu'),
10    layers.Dense(10, activation='softmax')
11 ])
```

[34]

Python

COMPILATION OF THE MODEL

```
1 cnn.compile(optimizer='adam',
2             loss='sparse_categorical_crossentropy',
3             metrics=['accuracy'])
```

[35]

Python

TRAINING OF THE CNN MODEL

```
1 cnn.fit(x_train, y_train, epochs=10) 🚀
```

[36]

Python

```
... Epoch 1/10  
1563/1563 [=====] - 40s 25ms/step - loss: 1.4798 - accuracy: 0.4674  
Epoch 2/10  
1563/1563 [=====] - 34s 22ms/step - loss: 1.1286 - accuracy: 0.6028  
Epoch 3/10  
1563/1563 [=====] - 35s 22ms/step - loss: 0.9963 - accuracy: 0.6513  
Epoch 4/10  
1563/1563 [=====] - 36s 23ms/step - loss: 0.9148 - accuracy: 0.6793  
Epoch 5/10  
1563/1563 [=====] - 36s 23ms/step - loss: 0.8516 - accuracy: 0.7027  
Epoch 6/10  
1563/1563 [=====] - 35s 23ms/step - loss: 0.7992 - accuracy: 0.7240  
Epoch 7/10  
1563/1563 [=====] - 36s 23ms/step - loss: 0.7459 - accuracy: 0.7404  
Epoch 8/10  
1563/1563 [=====] - 35s 23ms/step - loss: 0.7022 - accuracy: 0.7551  
Epoch 9/10  
1563/1563 [=====] - 34s 22ms/step - loss: 0.6607 - accuracy: 0.7694  
Epoch 10/10  
1563/1563 [=====] - 36s 23ms/step - loss: 0.6208 - accuracy: 0.7844  
  
<keras.callbacks.History at 0x19a56155940>
```

```
1 from keras.models import load_model
2 loaded_model = load_model("annproject1.h5")
```

[13] ✓ 4.5s

Python

```
1 from sklearn.metrics import confusion_matrix , classification_report
2 import numpy as np
3 y_pred = loaded_model.predict(x_test)
4 y_pred_classes = [np.argmax(element) for element in y_pred]
5
6 print("Classification Report: \n", classification_report(y_test, y_pred_classes))
```

[13]

Python

... 313/313 [=====] - 4s 7ms/step

Classification Report:

	precision	recall	f1-score	support
0	0.68	0.81	0.74	1000
1	0.76	0.83	0.79	1000
2	0.68	0.49	0.57	1000
3	0.53	0.51	0.52	1000
4	0.67	0.60	0.64	1000
5	0.60	0.63	0.61	1000
6	0.72	0.81	0.76	1000
7	0.74	0.77	0.75	1000
8	0.82	0.78	0.80	1000
9	0.79	0.73	0.76	1000
accuracy			0.70	10000
macro avg	0.70	0.70	0.69	10000
weighted avg	0.70	0.70	0.69	10000

```
1 loaded_model.evaluate(x_test,y_test)
```

[14]

Python

```
... 313/313 [=====] - 5s 14ms/step - loss: 0.9368 - accuracy: 0.6971  
  
[0.9367502927780151, 0.6970999836921692]
```

```
1 y_pred = loaded_model.predict(x_test)  
2 y_pred[:5]
```

[15]

Python

```
... 313/313 [=====] - 3s 11ms/step  
  
array([[5.3353757e-05, 3.4360433e-05, 2.4336561e-05, 9.3526763e-01,  
        3.9762934e-05, 6.2665984e-02, 4.0117319e-04, 1.0616613e-04,  
        1.4036268e-03, 3.6480708e-06],  
       [1.2870802e-03, 5.9926391e-01, 5.6438996e-05, 5.2230496e-07,  
        1.8217817e-07, 1.1667767e-07, 4.4908387e-07, 1.5154913e-07,  
        3.9790127e-01, 1.4897849e-03],  
       [1.0875955e-01, 7.1699786e-01, 1.1480342e-03, 3.4640683e-03,  
        6.1841414e-04, 8.2107406e-04, 7.4084924e-04, 2.5710468e-03,  
        1.1531218e-01, 4.9566925e-02],  
       [9.0734780e-01, 6.4525456e-04, 1.2378589e-02, 1.0148327e-02,  
        3.8362953e-03, 2.5029414e-04, 4.5067168e-04, 4.3927989e-04,  
        6.4393640e-02, 1.0987990e-04],  
       [1.3995085e-05, 1.0697803e-06, 2.7654157e-03, 4.9842246e-02,  
        5.9531069e-01, 1.6057132e-02, 3.3562270e-01, 1.6582448e-05,  
        3.7022788e-04, 2.8315108e-08]], dtype=float32)
```




```
1 loaded_model.summary()
```

[14] ✓ 0.1s

Python

... Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 30, 30, 32)	896
max_pooling2d (MaxPooling2D)	(None, 15, 15, 32)	0
conv2d_1 (Conv2D)	(None, 13, 13, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 6, 6, 64)	0
flatten (Flatten)	(None, 2304)	0
dense (Dense)	(None, 64)	147520
dense_1 (Dense)	(None, 10)	650

=====

Total params: 167,562

Trainable params: 167,562

Non-trainable params: 0

=====

OUTPUT:

```
1 y_classes = [np.argmax(element) for element in y_pred]
2
```

[19]


Python

```
1 plot_sample(x_test, y_test, 66)
```

[25]

Python

...



```
1
2 classes[y_classes[66]]
```

[26]

Python

... 'automobile'